

# Databases

## Data-Modelling with ER

COMP 1531, 17s2

Aarthi Natarajan

Week 9

# Story so far...

Last week, we studied

- **relational data model**
- how to map the relational model to define the **database schema** (e.g., for a relational data model, define the tables, rows and columns or field names and types of fields, constraints and relationships)

This week, we will look at:

- **ER model**
- Mapping an ER model **to** a relational model

# Entity-Relationship (ER) Model

# Designing a database

**Data modelling:** an important early stage of database application development (aka "database engineering")

Typical steps in a database design

1. requirements analysis (identify data and operations)
- 2. data modelling** (high-level, abstract)
3. database schema design (detailed, relational model/tables)
4. database implementation (create instance of schema)
5. build operations/interface (SQL, stored procedures, GUI)
6. performance tuning (physical re-design)
7. schema evolution (logical schema re-design)


# Designing a database

## Two data models

- **Logical**: abstract model e.g., ER Model, OO Model
- **Physical**: record-based models e.g., relational model

## A strategy for designing a database

- Design using abstract model (conceptual-level modelling)
- Map to physical model (implementation-level modelling)



# Some Design Ideas

Consider the following while we work through exercises:

- start simple ... evolve design as problem better understood
- identify objects (and their properties), then relationships
- most designs involve kinds (classes) of people
- keywords in requirements suggest data/relationships  
(rule-of-thumb: nouns → data, verbs → relationships)
- don't confuse operations with relationships  
(**operation**: he **buys** a book; **relationship**: the book **is owned** by him)
- consider all possible data, not just what's available

# Quality of Designs

- There is no single "best" design for a given application
- Most important aspects of a design (data model):
  - correctness (satisfies requirements accurately)
  - completeness (all reqs covered, all assumptions explicit)
  - consistency (no contradictory statements)
- Potential inadequacies in a design:
  - omits information that needs to be included
  - contains redundant information ( $\Rightarrow$  inconsistency)
  - leads to an inefficient implementation
  - violates syntactic or semantic rules of data model

# Entity-Relationship Data Modelling

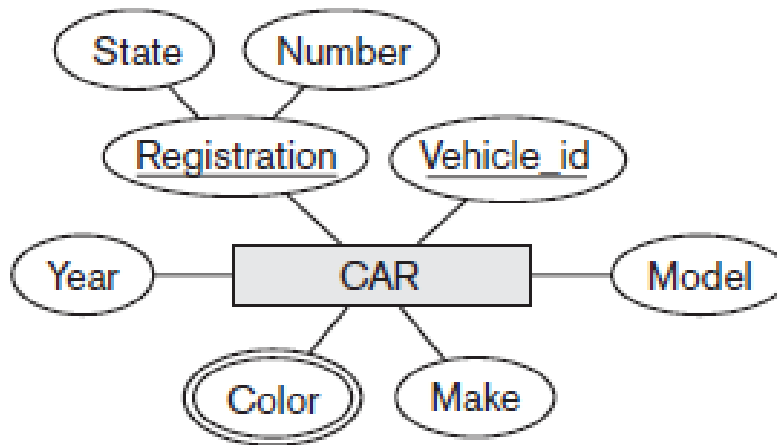
The world is viewed as a collection of **inter-related** entities.

ER modelling uses **three** major modelling constructs:

- **entity**:
  - a thing or object of interest in the real-world and is distinguishable from other objects
- **attribute**:
  - a data item or property of interest describing the entity  
*e.g., Joe (entity) described by name, address, age (attributes)*
- An **entity-set** (aka: entity-type) can be viewed as either:
  - a set of entities with the same set of attributes
  - an abstract description of a class of entities e.g., students, courses, accounts



e.g.,



An **entity-set** CAR with two **key attributes** (registration and vehicle\_id), three **single-valued attributes** (year, model, make) and a **multi-valued attribute** (color)

CAR  
Registration (Number, State), Vehicle\_id, Make, Model, Year, {Color}

CAR<sub>1</sub>  
((ABC 123, TEXAS), TK629, Ford Mustang, convertible, 2004 {red, black})

CAR<sub>2</sub>  
((ABC 123, NEW YORK), WP9872, Nissan Maxima, 4-door, 2005, {blue})

CAR<sub>3</sub>  
((VSY 720, TEXAS), TD729, Chrysler LeBaron, 4-door, 2002, {white, blue})

⋮

An **entity set** CAR with three **entities**

# Relationship Sets

**Relationship**: relates two or more entities

- e.g Joe Smith (entity) *is enrolled in* (relationship) COMP1531 (entity)

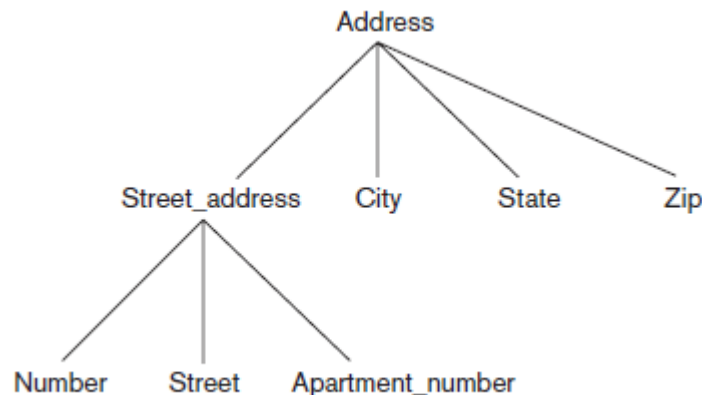
**Relationship Set (aka relationship type)** : set of similar relationships

- **degree** = # entities involved in the relationship (in ER model,  $\geq 2$ )
- **cardinality** = # associated entities on each side of relationship

# Attributes of an entity-set

In contrast to relational model, attributes in an ER model can be:

- **Simple** (attribute cannot be broken into smaller sub-parts)
  - e.g., **age** attribute for entity type Employee
- **Composite** (have a hierarchy of attributes)
  - e.g., entity type EMPLOYEE has a composite attribute **Address**



- **Single-valued** (have only one value for each entity)
  - e.g., an **vin\_chassis** attribute for an entity type CAR
- **Multi-valued** ( have a set of values for each entity)
  - e.g., a **Colors** attribute for CAR = (blue,black)

# What if two entities have the same set of attribute values?

- They're regarded as the same entity.
- So, each entity must have a distinct set of attribute values.

## One approach:

Define a **key (superkey)** : It is any set of attributes

- whose set of values are distinct over entity set
- natural (e.g. **name + address + birthday**) or artificial (e.g. **SSN**)

- **Candidate key** = any superkey such that **no subset** is also a superkey)  
e.g. (**name + address**) is a superkey, but not (**name**) or (**address**)
- **Primary key** = a candidate key chose by DB designer that uniquely identifies an entity e.g., **SSN**

## Example (bank customer entities)

Customer = (custNo, name, address, taxFileNo)

- Definite superkeys:
  - any set of attributes involving custNo or taxFileNo
- Possible superkeys:
  - (name,address)
- Unlikely superkeys:
  - (name), (address)

# ER model vs OO model

Analogy between ER and OO models:

- an **entity** is like an **object instance**
- an **entity set** is like a **class**

Differences between ER and OO models:

- ER modelling doesn't consider operations (methods)

# Entity Relationship Diagrams

- **ER diagrams** are a graphical tool for data modelling
- An ER diagram consists of:
  - a collection of *entity set* definitions
  - a collection of *relationship set* definitions
  - *attributes* associated with entity and relationship sets
  - connections between entity and relationship sets

**Warning:** 99% of the time ...

- we say "entity" when we mean "entity set"
- we say "relationship" when we mean "relationship set"
- If we want to refer to a specific entity, we generally say "entity instance"

# Entity Relationship Diagrams


Specific visual symbols indicate different ER design elements:




Entity



Weak entity




Relationship



Identifying Relationship



Attribute



Multi-valued attribute



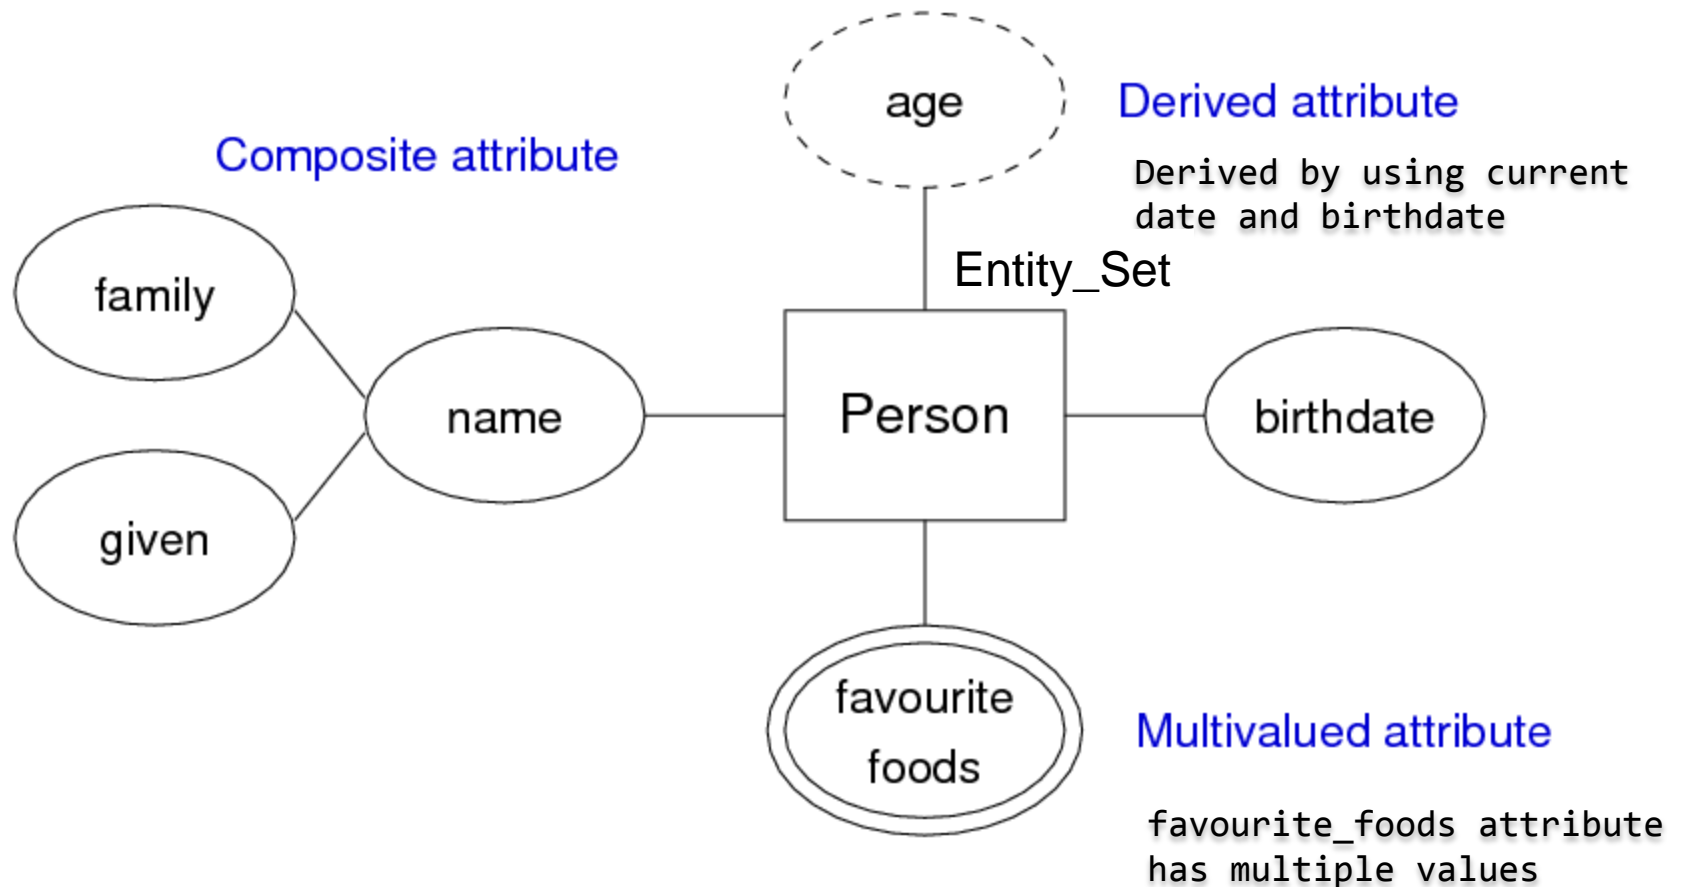
Inheritance



Derived attribute



# Example of attribute notations




# Cardinality in Relationship Sets

## Examples:


one-to-one



one-to-many



many-to-many



# An alternative explicit notation

## Examples:

one-to-one



one-to-many



many-to-many



# Relationship Sets in ER diagrams

**Level of participation constraint** = a type of relationship constraint defined as:

Participation in relationship set  $R$  by entity set  $A$  may be:

- **total** - every  $a \in A$  participates in  $\geq 1$  relationship in  $R$
- **partial** - only some  $a \in A$  participate in relationships in  $R$

Example:

- every bank loan is associated with at least one customer
- not every customer in a bank has a loan



# Exercise 1: Relationship Semantics

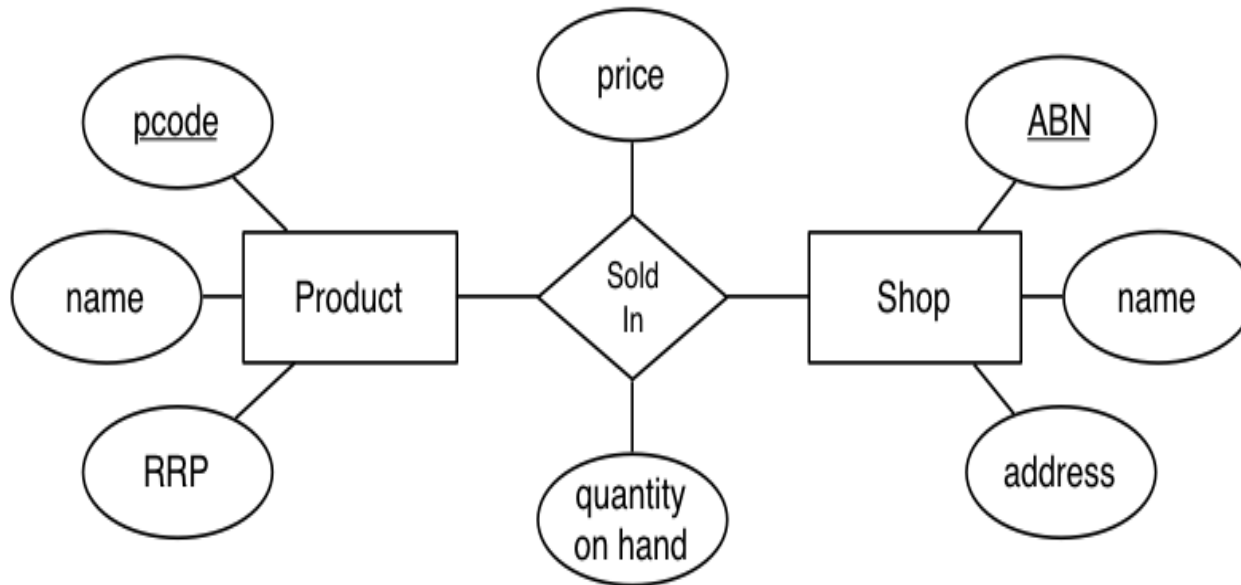
Describe precisely the scenarios implied by the following relationships:



# Relationship Type with attributes

In some cases, a relationship needs associated attributes

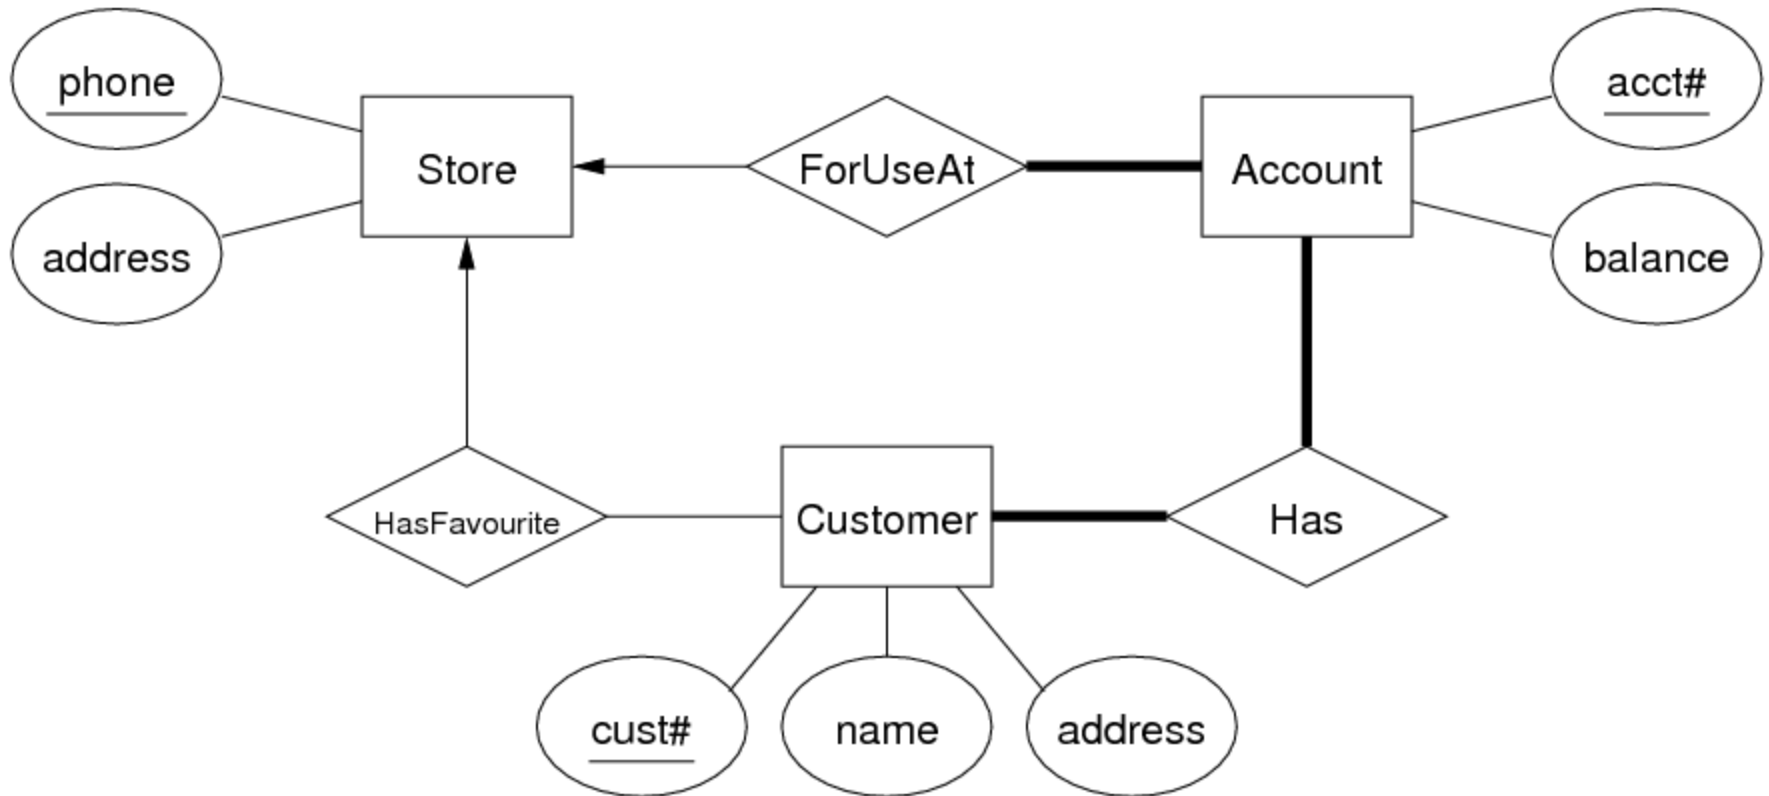
**Example:**



(price and quantity are related to products in a particular shop)

# Putting it all together...

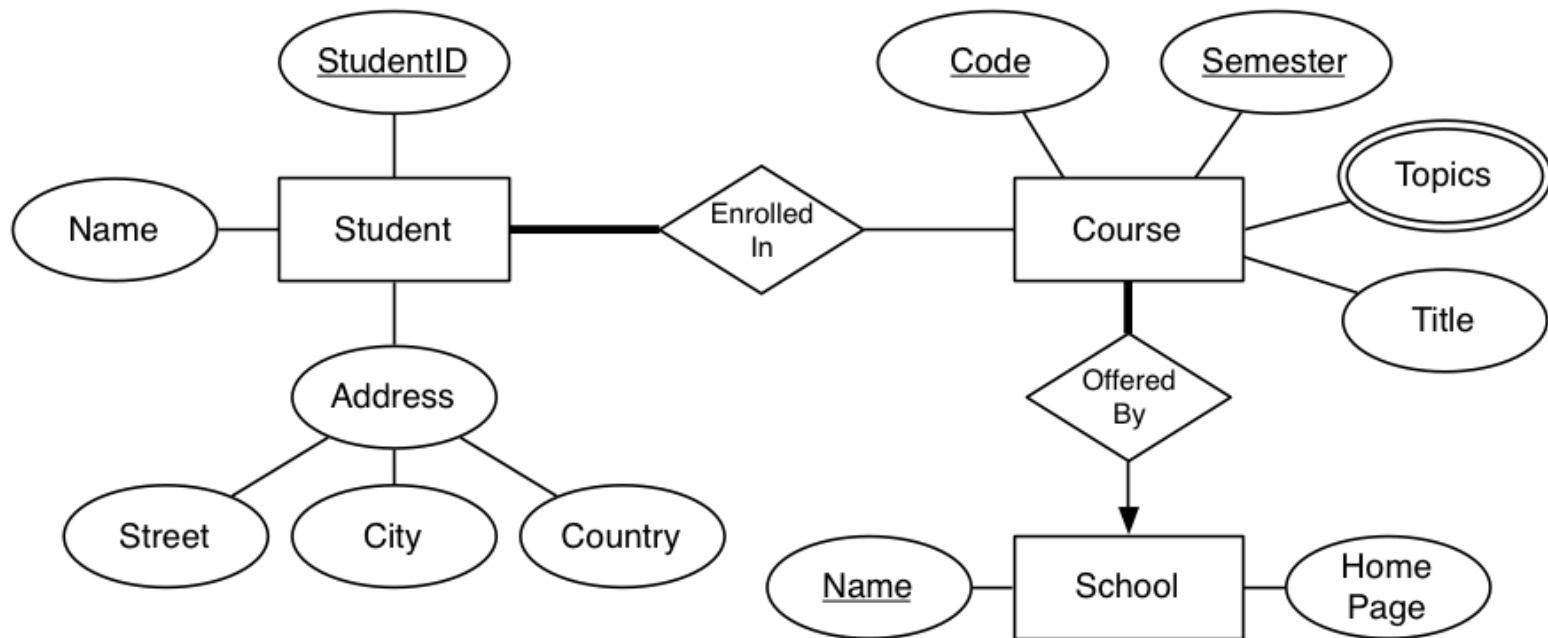
## Example1: - a complete ER Diagram



primary key attributes are underlined e.g. cust#

## Example 2:

Entities, relationships, attributes, keys, cardinality, participation, ...






# Weak Entity Set

A Weak entity set

- has no key of its own;
- exist only because of association with strong entities

Example:



# Subclasses and Inheritance

A **subclass** of an entity set  $A$  is a set of entities:

- with all attributes of  $A$ , plus (usually) it own attributes
- that is involved in all of  $A$ 's relationships, plus its own

Properties of subclasses:


- **overlapping** or **disjoint** (can an entity be in multiple subclasses?)
- **total** or **partial** (does every entity have to also be in a subclass?)

Special case: entity has one subclass ("B is-a A" specialisation)


# Subclasses and Inheritance

## Example:

*A person may be a doctor and/or may be a patient or may be neither*



*Every employee is either a permanent employee or works under a contract*




# Design considerations using the ER model

- should an "object" be represented by an attribute or entity?
- is a "concept" best expressed as an entity or relationship?
- should we use  $n$ -way rel<sup>n</sup>ship or several 2-way rel<sup>n</sup>ships?
- is an "object" a strong or weak entity? (usually strong)
- are there subclasses/superclasses within the entities?

Answers to above are worked out by *thinking* about the application domain.

# Design considerations (cont...)


## Attribute vs Entity Example (v1)



Employees can work for several departments, but cannot work for the same department over two different time periods.

# Design considerations (cont...)

## Attribute vs Entity Example (v2)



Employees can work for the same department over two different time periods.

# Design using the ER model


ER diagrams are typically too large to fit on a single screen.  
(or a single sheet of paper, if printing)

One commonly used strategy:

- define entity sets separately, showing attributes
- combine entities and relationships on a single diagram (but without showing entity attributes)
- if very large design, may use several linked diagrams as seen in the example in the next three set of slides


# e.g. an ER model for a Bank

## (1) Modelling people (employees)



## (2) Modelling people (customer)

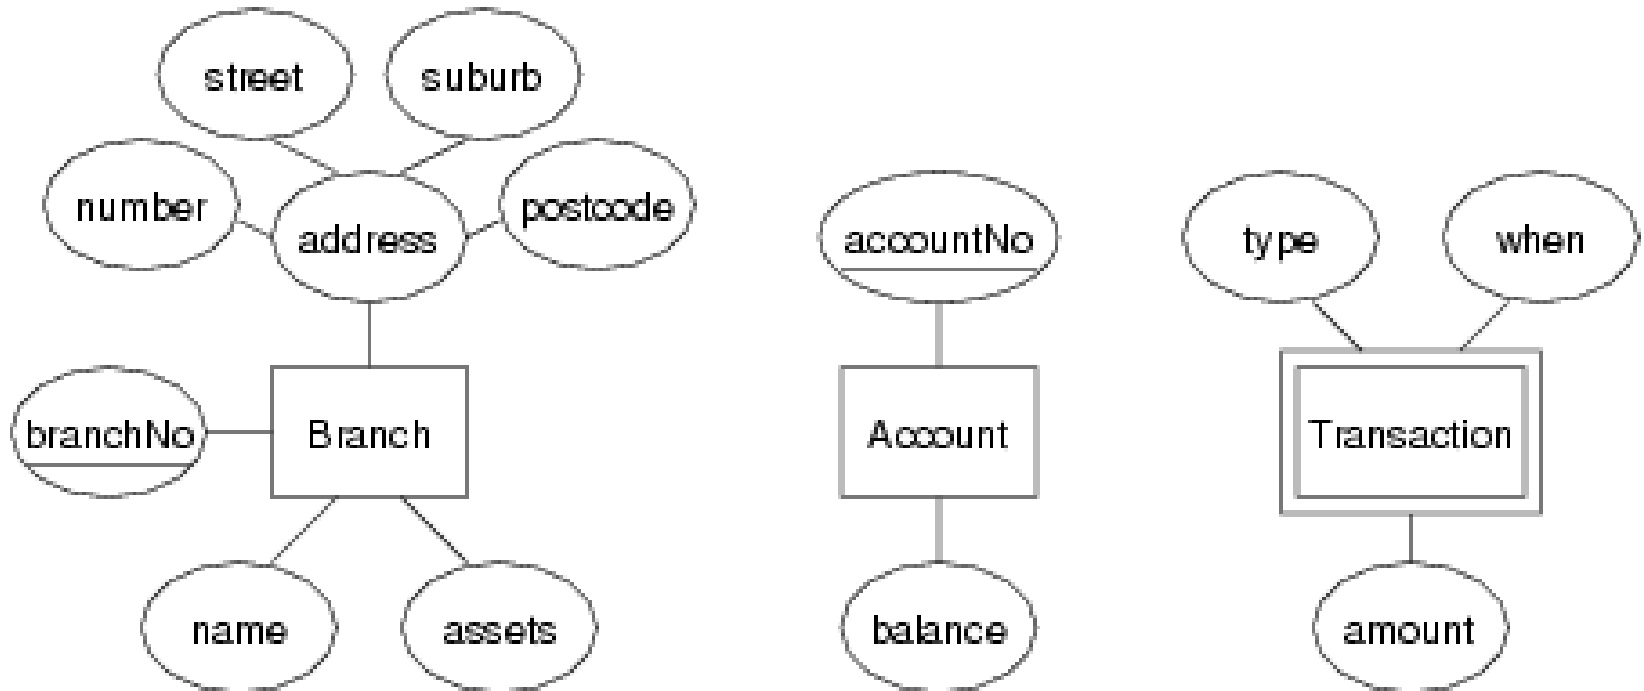
Modelling people (cont):






## e.g. an ER model for a Bank

### (3) Modelling branches, accounts, transactions



# e.g. an ER model for a Bank

## (4) Putting it all together with relationships



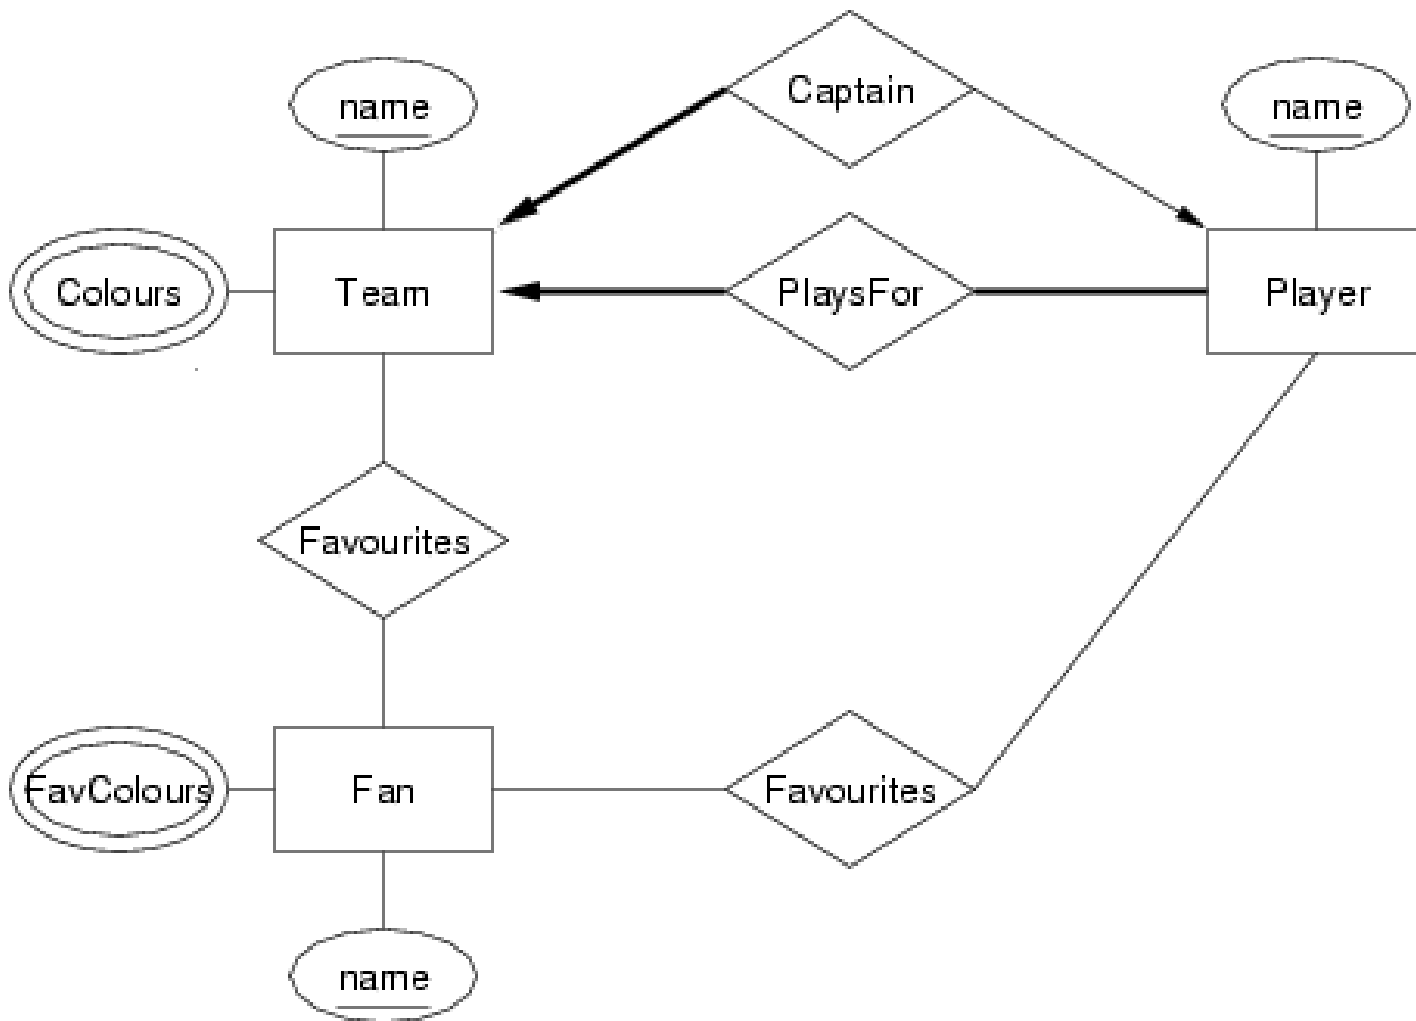
## Exercise 2:

Develop an ER design for the following scenario:

A database records information about teams, players, and their fans, including:

- For each team, its name, its players, its captain (one of its players) and the colours of its uniform.
- For each player, their name and team.
- For each fan, their name, favourite teams, favourite players, and favourite colour.

## Exercise 2 (Solution): ER Design



# Mapping ER Designs to Relational Model

## ER to Relational Mapping

- A formal mapping exists for ER model  $\rightarrow$  Relational model.
- This maps "structures"; but additional info is needed, e.g.
  - concrete domains for attributes and other constraints

# Relational Model vs Entity Model

Correspondences between relational (R) and ER data models:

- ER **attribute** → relational **attribute**
- ER **entity** → relational **tuple**
- ER **entity-set** → relational **table** (relation)
- ER **relationship** → relational **table** (relation)
- ER **key** → relational **primary key**

Differences between relational and ER models:

- Relational uses *relations* to model *entities* and *relationships*
- Relational has **no** *composite* or *multi-valued* attributes (only atomic)
- Relational has **no** *object-oriented notions* (e.g. subclasses, inheritance)

# (1) Mapping Strong Entities

An *entity* consists of:

- a collection of attributes;
  - attributes are simple, composite and multi-valued

A *relation schema* consists of:

- a collection of attributes;
  - all attributes have **atomic** data values

So, even the **mapping** from entity to relation schema is **not simple**




# (1) Mapping Strong Entities

An obvious mapping

- an entity set  $E$  with **atomic attributes**  $a_1, a_2, \dots, A_n$   
**maps to**
- a relation (table)  $R$  with **attributes (columns)**  $a_1, a_2, \dots, A_n$

Each row in relation  $R$  corresponds to an entity in  $E$

**Example:**



(Note: the key is preserved in the mapping)


## (2) Mapping Relationships

ER **relationship** → relational **table** (relation)

- Identify one entity as “parent”
- other entity as “child”
- as general rule,
  - **PK of parent is added to child as FK**
- Any attributes of the relationship
  - are added to **child** relation

## (2a) Mapping N:M Relationships

A binary relationship set  $B$  between entity sets  $E$  and  $F$  gives **associations** between pairs of entities in  $E$  and  $F$



We can represent

- entity set  $E$  by relation  $S$  (using attribute mappings as above)
- entity set  $F$  by relation  $T$  (using attribute mappings as above)

But how to represent  $B$ ?


## One possibility –

- Represent the **relationship** set  $B$  explicitly by a **relation**  $R$  containing:
  - all attributes from the primary keys of  $S$  and  $T$
  - all attributes associated with the relationship set  $B$
- where  $S$  and  $T$  are relations representing entity sets  $E$  and  $F$
- and the **key** for  $R$  is the **union of the key attributes for  $S$  and  $T$**

And this approach works generally for:

- relationship degree  $\geq 2$
- relationship multiplicity **1:1**, **1:N**, **N:M**
- associated attributes are simply included in  $R$

# Example - Mapping N:M Relationship



## *Relational Version*

Customer

<u>custNo</u>	name	address
---------------	------	---------

Account


<u>acctNo</u>	branch	balance
---------------	--------	---------

Owns

<u>custNo</u>	<u>acctNo</u>	opened
---------------	---------------	--------

## (2b) Mapping 1:M Relationships

### Example:



### Relational Version

#### Generic Mapping

Customer

<u>custNo</u>	name	address
---------------	------	---------

Branch

<u>branchNo</u>	address	assets
-----------------	---------	--------

Home

<u>custNo</u>	<u>branchNo</u>	joined
---------------	-----------------	--------

#### Optimised Mapping

Customer


<u>custNo</u>	name	address	branchNo	joined
---------------	------	---------	----------	--------

Branch

<u>branchNo</u>	address	assets
-----------------	---------	--------

## (2c) Mapping 1:1 Relationships

### Example:

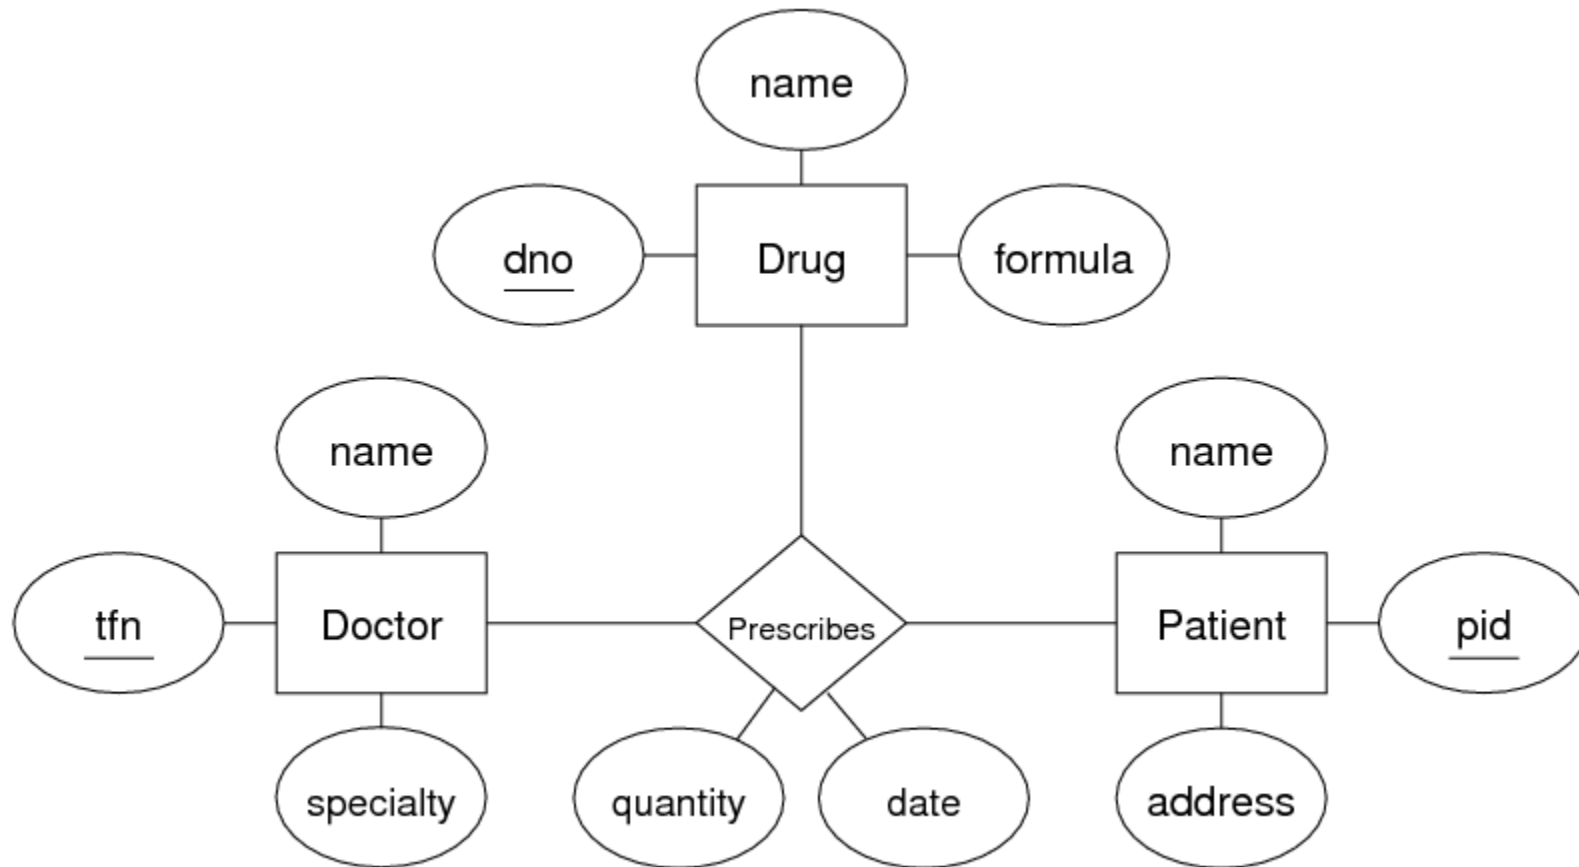


- Handled similarly to 1:N relationships
- For a 1:1 relationship between entity sets  $E$  and  $F$  ( $S$  and  $T$ ):
  - choose one of  $S$  and  $T$  (e.g.  $S$ ) (*Note : Choose the entity set that participates totally, if only one of them does*)
  - add the attributes of  $T$ 's primary key to  $S$  as foreign key
  - add the relationship attributes as attributes of  $S$

## (2d) Mapping n-way relationships

### Exercise:


Convert the following ER design into a relational data model





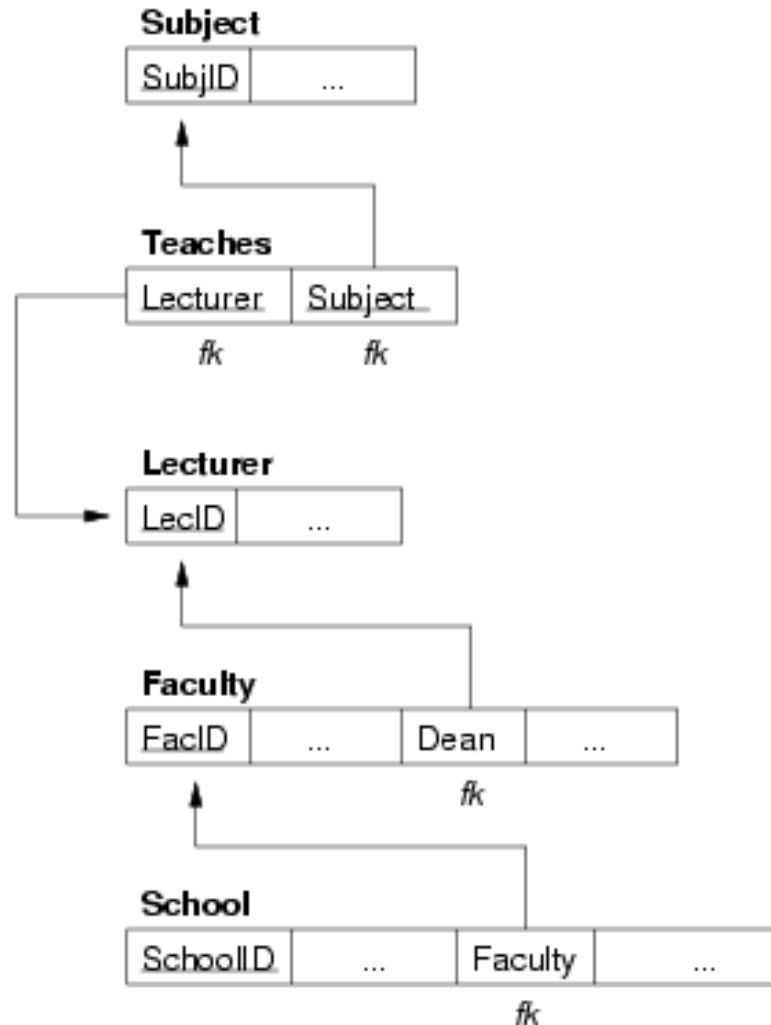
## Exercise3:

Convert the following ER design into a relational data model



*You can assume that each attribute contains (at least) a suitably-named attribute containing a unique identifying number (e.g. the Lecturer entity contains a LecID attribute).*

## Exercise3( Solution ) : Relational Model



Relational model for a very small University ER model

**Wednesday**

# So far ...

## Summarising ER → Relational Mapping

### Mapping entities and attributes

- ER **attribute** → relational **attribute**
- ER **entity** → relational **tuple**
- ER **entity-set** → relational **table** (relation)
- ER **key** → relational **primary key**

### Mapping Relationships


- ER **relationship** → create a **new** relational **table** (relation)
  - N:M relationship → add FK for each participating entity plus relationship attributes)
  - 1:N relationship → FK plus relationship attributes
  - 1:1 relationship → FK plus relationship attributes

### (3) Mapping multi-valued attributes

- treat like an **N:M relationship** between **entities** and **values**
- create a **new relation** where each tuple contains:
  - the primary key attributes from the entity
  - one value for the multi-valued attribute from the corresponding entity

**Example:**

*ER Model*



*Relational Version*

Person

<u>SSN</u>	name	birthdate
------------	------	-----------

FavColour

<u>SSN</u>	<u>colour</u>
------------	---------------

### (3) Mapping multi-valued attributes contd...

**Example:** the two entities

```
Person(12345, John, 12-feb-1990, [red, green, blue])  
Person(54321, Jane, 25-dec-1990, [green, purple])
```

would be represented as:

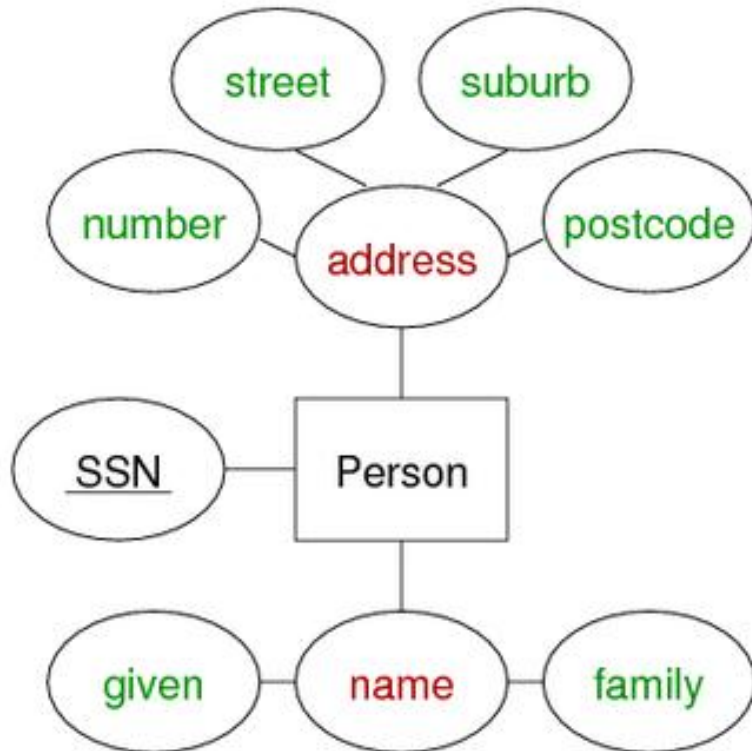
```
Person(12345, John, 12-feb-1990)  
Person(54321, Jane, 25-dec-1990)  
  
FavColour (12345, red)  
FavColour(12345, green)  
FavColour(12345, blue)  
FavColour(54321, green)  
FavColour(54321, purple)
```

## (4) Mapping composite attributes

Mapped by concatenation or flattening

Example:

*ER Model*



*Relational Version #1*

Person

<u>SSN</u>	name	address
------------	------	---------


*Relational Version #2*

Person

<u>SSN</u>	givenName	familyName	...	
...	number	street	suburb	postcode

## Exercise 4:

Convert the following ER design to relational form



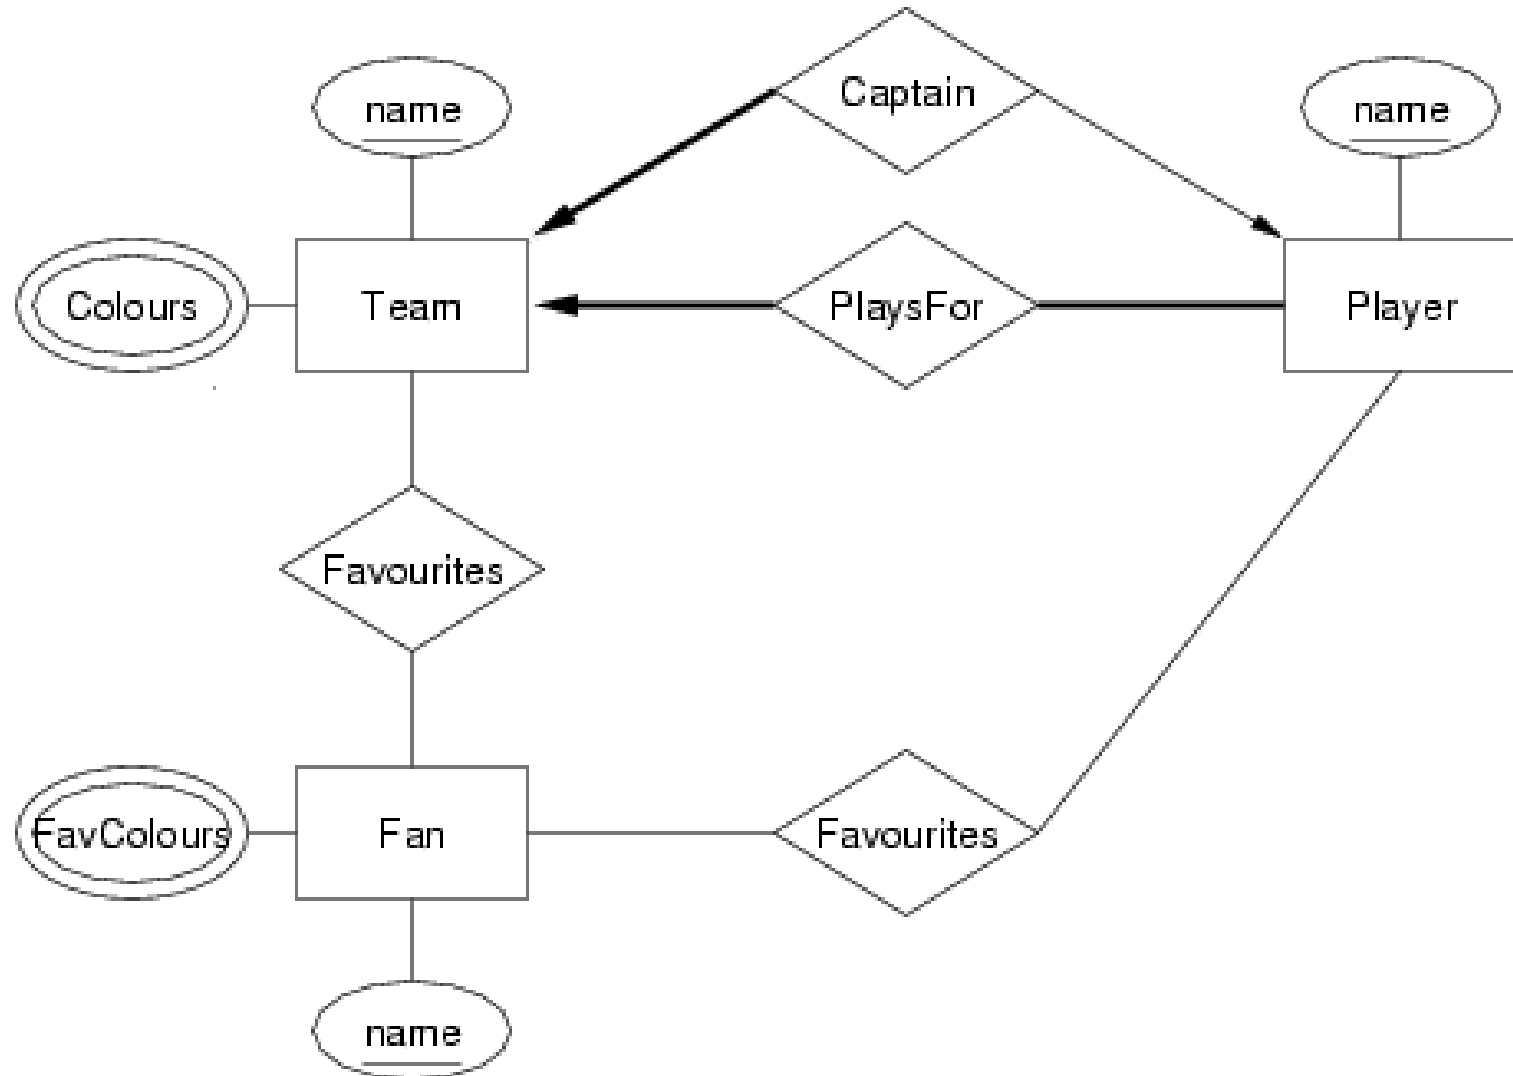


## Exercise 5:

Convert the following ER design into a relational data model

(1) first as a box-and-arrow diagram

(2) as a sequence of statements in the SQL data definition language:



## (5) Mapping sub classes

Three different approaches to mapping subclasses to tables:


- **ER style**
  - each entity becomes a separate table,
  - containing attributes of subclass + FK to superclass table
- **object-oriented**
  - each entity becomes a separate table,
  - inheriting all attributes from all superclasses
- **single table with nulls**
  - whole class hierarchy becomes one table,
  - containing all attributes of all subclasses (null, if unused)

*Which mapping is best depends on how data is to be used*

## (5) Mapping sub classes in ER style

The subclass relation contains:

- all of the subclass-specific attributes
- uses the superclass primary key to capture the association



*Relational Model*

Person

SSN	name	address
-----	------	---------

Employee


SSN	salary	position
-----	--------	----------

Manager

SSN	bonus
-----	-------

## (5) OO Mapping of sub classes

*Entity-Relationship Model*



*Relational Model*

Person

<u>SSN</u>	name	address
------------	------	---------

Employee


<u>SSN</u>	name	address	salary	position
------------	------	---------	--------	----------

Manager

<u>SSN</u>	name	address	salary	position	bonus
------------	------	---------	--------	----------	-------

## (5) Single table with nulls mapping

*Entity-Relationship Model*



*Relational Model*


Person

<u>SSN</u>	name	address	salary	position	bonus
------------	------	---------	--------	----------	-------

*NULL for Person who is not Employee*

*NULL for Employee who is not Manager*

## (5) Exercise 6: Disjoint subclasses



Use (a) ER-mapping, (b) OO-mapping, (c) 1-table-mapping  
Are there aspects of the ER design that can't be mapped?

## Exercise 7: Give an ER design to model the following scenario ...

- Patients are identified by an **SSN**, and their names, addresses and ages must be recorded.
- Doctors are identified by an SSN. For each doctor, the name, specialty and years of experience must be recorded.
- Each pharmacy has a name, address and phone number. A pharmacy **must** have a manager.
- A pharmacist is identified by an SSN, he/she can only work for one pharmacy. For each pharmacist, the name, qualification must be recorded.
- For each drug, the trade name and formula must be recorded.
- Every patient has a primary physician. Every doctor has at least one patient.
- Each pharmacy **sells** several drugs, and has a price for each. A drug could be sold at several pharmacies, and the price could vary between pharmacies.
- Doctors prescribe drugs for patients. A doctor could prescribe one or more drugs for several patients, and a patient could obtain prescriptions from several doctors. Each prescription has a date and quantity associated with it.